

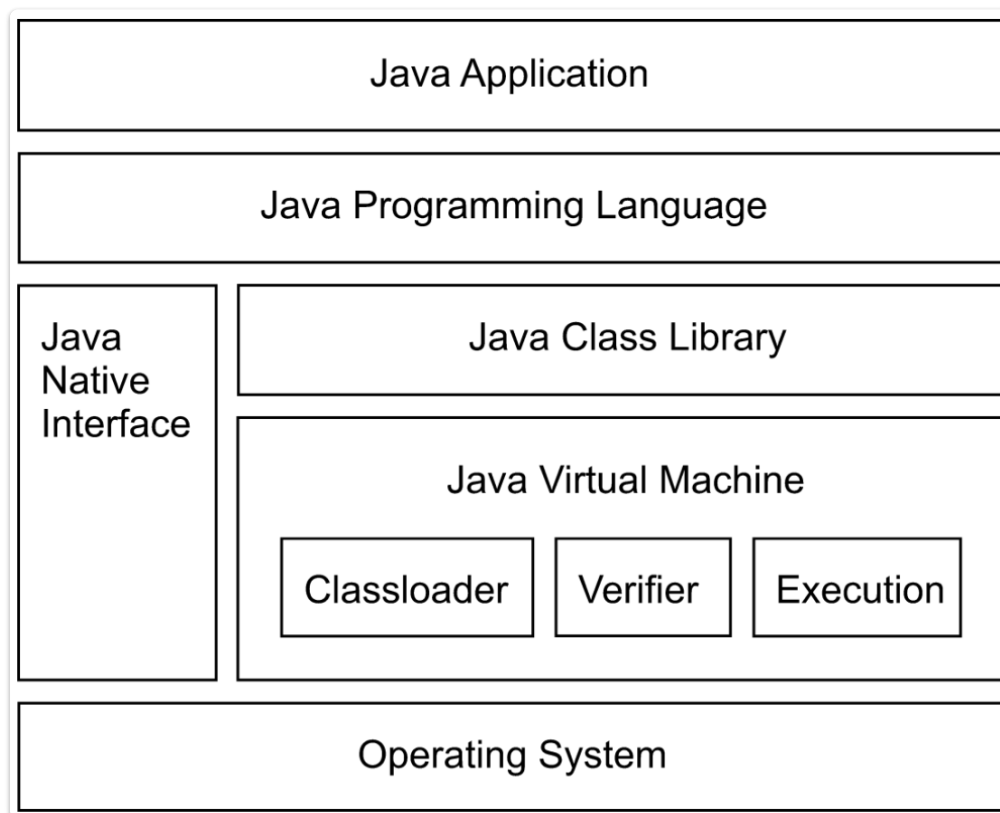
03 - Java Virtual Machine

Is a **multi-threaded stack based machine**: each instruction **take** arguments from **top** of the operand stack and also **put** results on **top** of operand stack (**LIFO**).

It defines the results of the **compilation** in **bytecode** (JVM's language) as a **machine independent class file format** (**.class**), that must be supported by all JVM implementations. It is **abstract**, so no implementation details like memory layout of run-time data area, garbage-collection, optimizations... Also **object representation** is left to the implementation, even the **null** value.

Note that **.class** is platform independent, but the loading process transforms **.class** to an **internal representation**, which is **implementation dependent**.

Java Hierarchy



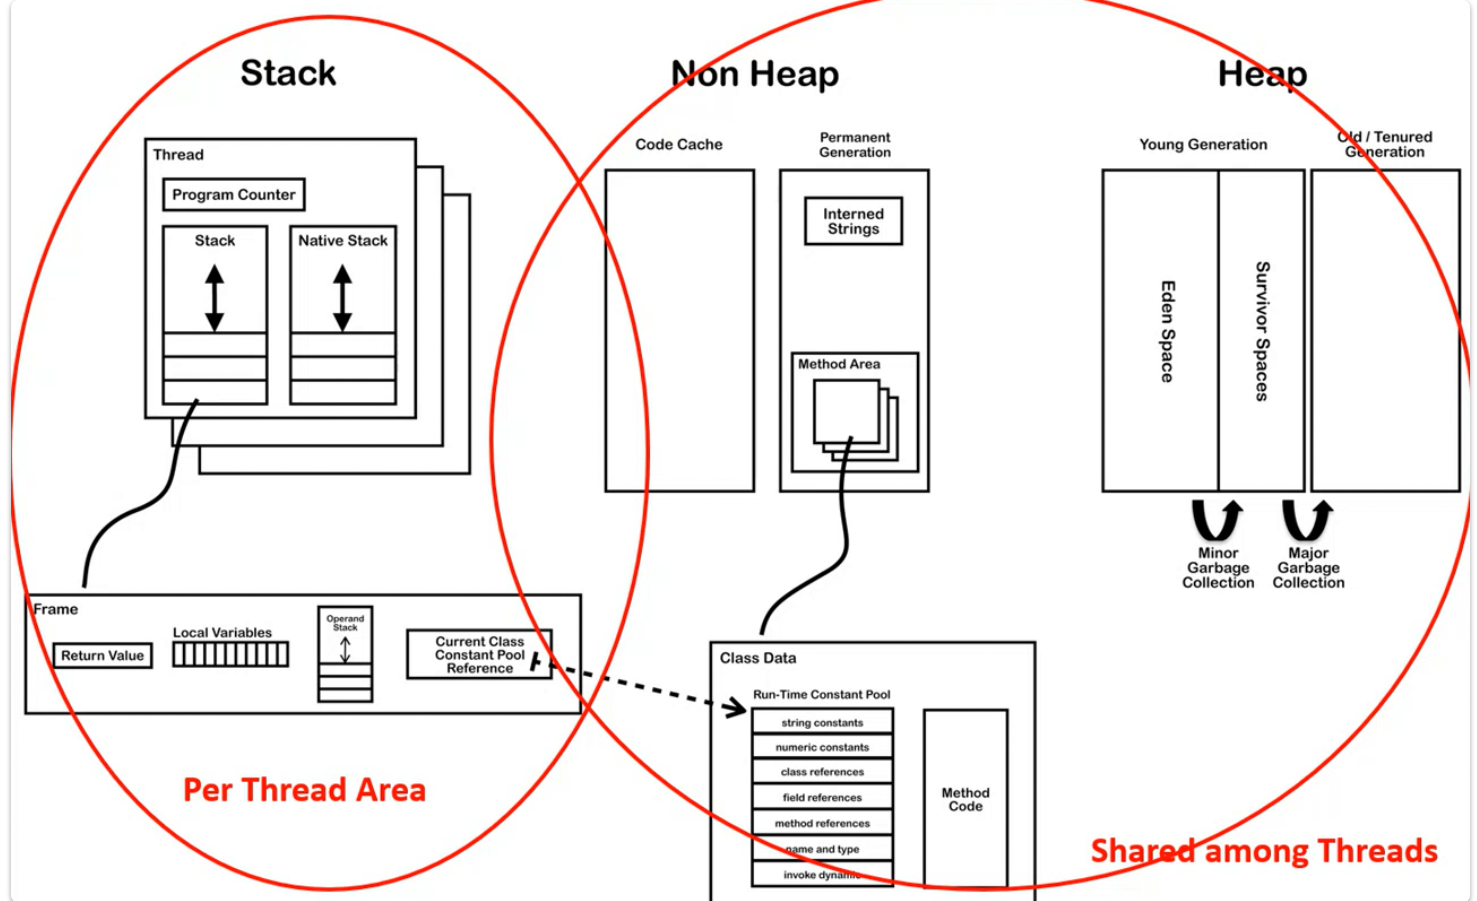
Thread

Multiple threads allowed per application (e.g. `main()` method), managed by **time-slicing** or **parallelization**. There is a **direct mapping** between a **Java Thread** and a **native** operating system Thread.

Instance of class `Thread`, **started** invoking `start()`. **States**: ready, in execution, suspended (`Thread.sleep()`), blocked (I/O or monitor request), waiting, terminated.

Daemon: **background system** threads (e.g. garbage collection, finalization, signal dispatching, ecc...).

Data Areas




- **Per Thread:**
 - **PC** (Program Counter);
 - **Java stack** (activation records -> method invocation/completion);
 - **Native stack** (Java Native Interface -> e.g. C function is invoked). Note that the frames in the stacks are not necessarily allocated contiguously, it could be implemented as a linked list (remember that JVM is an abstract machine);
- **Shared among Threads** (need **thread safety**):
 - **Heap**: objects and arrays (unlike C++). There is no explicit deallocation, the content will be garbage collected;
 - **Non Heap**: static part of the heap not garbage collected. Contains:
 - **Code cache** (JIT compilation): since interpreting **bytecode** is **not** as **fast** as native code, JVM look for bytecode executed **several time** and **compile** it to **native code**, stored in code cache;
 - **Permanent generation** (objects never deallocated):
 - **Interned strings** (string interning is a method of storing only **one copy** of each distinct string value, which must be **immutable**);
 - **Method Area** (memory where class files are loaded);

Example of class file:

ClassFile {	
u4 magic;	0xCAFEBABE
u2 minor_version;	Java Language Version
u2 major_version;	
u2 constant_pool_count;	Constant Pool
cp_info contant_pool[constant_pool_count-1];	
u2 access_flags;	access modifiers and other info
u2 this_class;	References to Class and Superclass
u2 super_class;	
u2 interfaces_count;	References to Direct Interfaces
u2 interfaces[interfaces_count];	
u2 fields_count;	Static and Instance Variables
field_info fields[fields_count];	
u2 methods_count;	Methods
method_info methods[methods_count];	
u2 attributes_count;	Other Info on the Class
attribute_info attributes[attributes_count];	
}	

Frame



- **Local Variables Array:**
 - Reference to **this**;
 - Method parameters;
 - Local variables.
- **Operand Stack:** expressions/methods evaluation;
- Reference to **Constant Pool**;
- **Return value**.

When a Java class is compiled, all the **constants** and all **references** to variables and methods are **stored** in the class **constant pool** as symbolic references. JVM can then resolve (**bind**) symbolic references in two ways:

- **Eager** (static): after class **file** is **loaded**;
- **Lazy** (late): when the **reference** is used for the **first time**.

This process is only done **once**, because the symbolic reference is completely replaced in the constant pool.

Method Area

It is the **memory** where **class files** are **loaded**, composed of:

- **ClassLoader Reference**;
- From the class file:
 - **Run Time Constant Pool**: constants / references (to fields/methods...);
 - **Field data**;
 - **Method data**;
 - **Method code** (bytecode);
 - ...

ClassLoader

All classes that are loaded contain a reference to the classloader that loaded them. In turn the classloader also contains a reference to all classes that it has loaded.

Generally works in **3 steps**:

- **Loading**: **read** the class file, **check** if it is properly formatted and all data are recognized by JVM and **create** an appropriate class/interface (from bytecode to the internal representation specific to the JVM implementation);
- **Linking**: take the previously created class/interface and **combine** it with the run-time state of the JVM so that it can be executed. It consists in 3 steps:
 - **Verification**: overflow/underflow, types validity...;
 - **Preparation**: allocation of storage (method tables);
 - **Resolution** (optional, because of eager/lazy binding): resolve symbol references by loading referred classes/interfaces.
- **Initialization**: **execute** the class/interface method `clinit()` (initialize class variables). Instead, when instances are initialized, the `init()` method is called.

We have **3 main types**:

- **Bootstrap Class loader**: written in **native code**, **loads basic Java APIs** (e.g.: `rt.jar`). During the JVM start-up, it **loads** the **initial class**;
- **Extension Classloader**: **middle-layer** loader, loads classes from standard **Java extension APIs** (e.g.: security extension functions);
- **System Classloader**: **default** application classloader, loads application classes from **classpath**.

Classloaders can also be **user-defined**, e.g.: runtime reloading of classes, loading from different sources (network, encrypted file...).

JVM Start-up and shutdown

This is the classic procedure when a **program starts**:

- The JVM starts up by **loading** an initial class (where the main method is) using the *bootstrap class loader*;
- The class is **linked** and **initialized**;
- The `public static void main(String[] args)` method is invoked;
- This will trigger **loading**, **linking** and **initialization** of additional classes and interfaces.

JVM **exits** when:

- all **non-daemon** threads **terminate**;
- `Runtime.exit`, `System.exit` (assuming it is secure).

Data Types

- **Primitives**: byte, short, int, long, char, float, double, boolean (supported only for arrays, since the smallest indexable word is of 1 byte and thus we can't read 1 bit, as the bool is), `returnAddress` (exception handling);
- **References**: class, array, interface;
- **Opcodes**: `iadd`, `fadd`... (types of operands).

However, there is **no type information** on **local variables at runtime**.

Instruction Set

Basic instructions that are executable on the virtual machine. They are the analogue of the instructions in machine language.

Properties:

- **Variable length**: **one-byte opcode** followed by **arguments**, against the stack machine of JVM (32 bit);
- Only **relative branches** (no GOTOs, jumps are done by only summing/subtracting values to the PC);
- **Byte aligned**: a data is x-byte aligned if it has a number of bytes which is multiple of x. The instruction set is 1-byte aligned, saving a lot of space but in order to read everything we must read 1 byte at a time;
- May contain **symbolic references**;

Examples:

- **Load** and **store** move back and forth data between the operand stack (where the stack machine JVM operates) and the local variable array;
- Arithmetic, control transfer, method invocation and return...

Moreover, each instruction could have **different forms** (like different constructors), e.g. different forms of "iload".

There are 3 addressing modes:

- **Intermediate**, a constant is part of instruction (e.g. `iload_1`: we know that we refer to the element in position 1 of the LVA);
- **Indexed**, access variable from local variable array;
- **Stack**, pop operand stack.

OpCodes are **explicitly typed**: the first letter of the OpCode identify the type. However, non-supported types will have to be converted, almost no support for **byte**, **char** and **short**. Instead, **int** is used as "computational type":

- i: int;
- l: long;
- s: short;
- b: byte;
- c: char;
- f: float;
- d: double;
- a: for reference.

Commands

- Compiler: `javac filename.java`;
- Disassembler: `javap -c -v filename.class` (option -c: produce entire bytecode, constant pool included);
- Execution: `java filename`.